

Course Review

CSCI3170 2024T1

Course – Learning Outcomes

The students will be able to

1. use an E-R diagram to model a database;
2. translate an E-R diagram into a relational model;
3. fine tune a relational schema based on the principles of relational database normalization;
4. implement queries by using database languages (SQL in particular);
5. understand file organizations and index structures of a DBMS;
- ~~6. understand the ideas of query processing and query optimization;~~
7. understand the principles of concurrency control and recovery schemes;

Overview: Database Design

- Data models: ER, Relational Data Model and their mapping
- Relational Algebra: Being able to use relational algebra to answer question.
- Relational Database Design: Functional Dependency, Normal Forms, Design Algorithms for 3rd normal form and B-C normal form (3.5 normal form)

Overview: RDBMS + Other DB


- Disk, Files, Buffer Replacement Policy
- Index: Introduction.
- Transaction Management
 - ACID properties
 - Conflict-serializable vs serializable
 - Concurrency control (locking, time-stamp ordering)

What's Important

- Why Database: Files vs. DBMS
- Database Modeling: Entity-relationship Model
- How it Works: Relational Algebra and SQL-based Application
- Make it Faster: Schema Refinement
- The Underlying: Storage vs. Memory, Indexes, and B+ Tree
- More: Hashing, Concurrency Control, and Recovery
- (Note: This short slide aims to help you organize what you've learned, please don't forget to review the more comprehensive lecture slide.)

Entity-Relationship diagram (E-R diagram)

- The E-R model can be presented graphically by an E-R diagram.



Relational data model

- Most DBMS today are based on the relational data model.
- Relation
 - the central data description construct in this model.
 - It can be thought of as a set of records.
 - A table with rows and columns.
 - Row – a record
 - Column – field, attribute.

Relational data model

| sid | name | login | age | gpa |
|-------|---------|---------------|-----|-----|
| 53666 | Jones | Jones@cs | 18 | 3.4 |
| 53688 | Smith | Smith@ee | 18 | 3.2 |
| 53650 | Smith | Smith@math | 19 | 3.8 |
| 53831 | Madayan | Madayan@music | 11 | 1.8 |
| 53832 | Guldu | guldu@music | 12 | 2.0 |

An example of a student relation

- A description of data in terms of a data model is called a schema.
- The schema of the above table is
Students(*sid*: string, *name*: string, *login*: string, *age*: integer, *gpa*: real)

Basic Operators

- There are six basic operators in relation algebra:
 - select (σ)
 - project (Π)
 - union (\cup)
 - set different ($-$)
 - Cartesian product (\times)
 - rename (ρ)

Motivation Example

- Redundant Storage
 - The rating value 8 corresponds to the hourly wage 10, and this association is repeated three times.
 - Storage is not used efficiently.

| <u>Id</u> | name | age | rating | Hourly_wages | Hours_worked |
|-------------|-------|-----|--------|--------------|--------------|
| 123-22-3666 | Peter | 48 | 8 | 10 | 40 |
| 231-31-5368 | Paul | 22 | 8 | 10 | 30 |
| 131-24-3650 | Mary | 35 | 5 | 7 | 30 |
| 434-26-3751 | David | 35 | 5 | 7 | 32 |
| 612-67-4134 | Ada | 35 | 8 | 10 | 40 |

- A decomposition of a relation schema R consists of replacing the relation schema by two (or more) relation schemas that each contains a subset of attributes of R and together include all attributes in R

| <u>Id</u> | name | age | rating | Hourly_wages | Hours_worked |
|-------------|-------|-----|--------|--------------|--------------|
| 123-22-3666 | Peter | 48 | 8 | 10 | 40 |
| 231-31-5368 | Paul | 22 | 8 | 10 | 30 |
| 131-24-3650 | Mary | 35 | 5 | 7 | 30 |
| 434-26-3751 | David | 35 | 5 | 7 | 32 |
| 612-67-4134 | Ada | 35 | 8 | 10 | 40 |

| <u>Id</u> | name | age | rating | Hours_worked |
|-------------|-------|-----|--------|--------------|
| 123-22-3666 | Peter | 48 | 8 | 40 |
| 231-31-5368 | Paul | 22 | 8 | 30 |
| 131-24-3650 | Mary | 35 | 5 | 30 |
| 434-26-3751 | David | 35 | 5 | 32 |
| 612-67-4134 | Ada | 35 | 8 | 40 |

| rating | Hourly_wages |
|--------|--------------|
| 8 | 10 |
| 5 | 7 |

Functional dependency:
- rating determines Hourly_wages

Disks and Files

- DBMS stores information on (“hard”) disks.
 - A disk is a sequence of bytes, each has a disk address.
 - READ: transfer data from disk to main memory (RAM).
 - WRITE: transfer data from RAM to disk.
 - Both are high-cost, relative to in-memory operations.
- Data are stored and retrieved in units called disk blocks or pages.
 - Each page has a fixed size, say 512 bytes. It contains a sequence of records.
 - In many (not always) cases, records in a page have the same size, say 100 bytes.
 - In many (not always) cases, records implement relational tuples.


Why Not Store Everything in Main Memory?

- Costs too much.
 - Ram is much more expensive than disk.
- Main memory is volatile.
 - We want data to be saved between runs. (Obviously!)
- Typical storage hierarchy:
 - Main memory (RAM) for currently used data. Disk for the main database (secondary storage).
 - Tapes for archiving older versions of the data (tertiary storage).

Indexes

- An index on a file speeds up selections on the search key fields
- for the index.
 - Any subset of the fields of a relation can be the search key for an index on the relation.
 - Search key is not the same as key (minimal set of fields that uniquely identify a record in a relation).
- An index contains a collection of data entries
 - A data entry is denoted as k^* , where k is a search key value and $*$ tells where to find the record containing k
 - Index must support efficient retrieval of all data entries k^* with a given key value k . Structure of data entry in more detail
- Primary vs. secondary: If search key contains primary key, then called primary index, otherwise secondary.
- Unique index: Search key contains a candidate key.

An example of a B+ tree

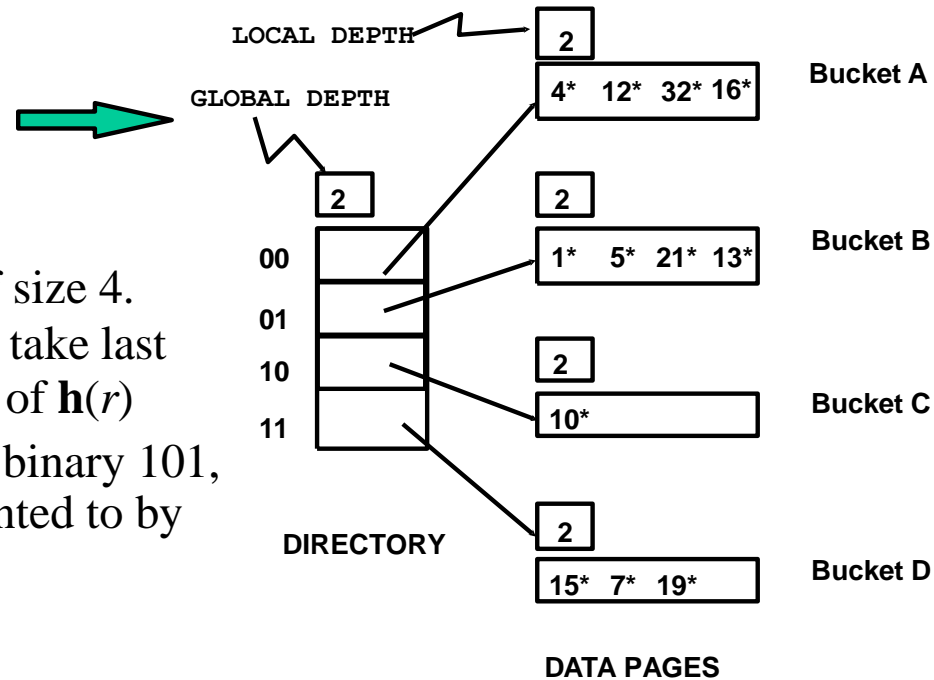


Extendible Hashing

- Situation: Bucket (primary page) becomes full. Why not re-organize file by doubling # of buckets?
 - Must re-hash all data entries to the right buckets
 - Example: assume hash function $h(k) = k \bmod M$
 - For $M = 4$, entries 3^* and 7^* both in bucket 3 ($3 \bmod 4 = 7 \bmod 4 = 3$)
 - But for $M = 8$, entry 7^* will be in bucket 7
 - Can we only re-hash those values that have changed addresses?
 - Difficulties: without re-hashing all the values, we don't know which values keep the old addresses and which get new addresses
 - Reading and writing all pages are expensive!
 - Question: how do we add more buckets, but only re-hash a few data entries?
 - Answer: use a level of indirection, directory of pointers to buckets

Example

- $h(r) = r \bmod 32$
- Directory is array of size 4.
- To find bucket for r , take last '*global depth*' # bits of $h(r)$
 - If $r = 5$, $h(r) = 5 = \text{binary } 101$, 5^* is in bucket pointed to by 01.



- ❖ **Insert:** If bucket is full, *split* it (allocate new page, re-distribute).
- ❖ *If necessary*, double the directory. (As we will see, splitting a bucket does not always require doubling; we can tell by comparing *global depth* with *local depth* for the split bucket.)

Concurrency

- Multiple users.
- Concurrent accesses.
- Problems could arise if there is no control.
- Example:
 - Transaction 1 (T1): withdraw \$500 from account A.
 - Transaction 2 (T2): deposit \$800 to account A.

```
Read(A)
A = A - 500
Write(A)
```

T₁

```
Read(A)
A = A + 800
Write(A)
```

T₂

Transactions

- A *transaction* is a sequence of read/write operations.
- A transaction is *atomic*.
 - Either all or none of the operations are executed.
 - No transaction can observe partial effect of other transactions.

Recoverable

- Transactions may be aborted due to logical failure. e.g. deadlock
- Recoverability is required to ensure that aborting a transaction does not change the semantics of committed transaction's operations.
- Example
 - $\text{Write}_1(x,2); \text{Read}_2(x); \text{Write}_2(y,3); \text{Commit}_2$
 - Not recoverable.
 - T_2 has committed before T_1 commits.
 - The problem is: what can we do if T_1 abort?
 - Delaying the commitment of T_2 can avoid this problem
- A schedule H is called *recoverable* (RC) if, whenever T_i reads from T_j . ($i \neq j$) in H and $c_i \in H$, $c_j < c_i$.
- Intuitively, a history is recoverable if each transaction commits after the commitment of all transactions (other than itself) from which it reads.

In Exam, I will ask about

- I will ask mostly about
 - ER Model, Relational Model
 - Relational Algebra
 - Storage, Indexes (*e.g. Buffers, B+ Trees*)
 - Transactions, Schedules, Concurrency
 - Recovery (**conceptual mostly**)
- None
 - View serializability
 - Differences on clustered or unclustered Indexes
 - Java
 - SQL Triggers/ Check Constraints

Final Exam/Assignment

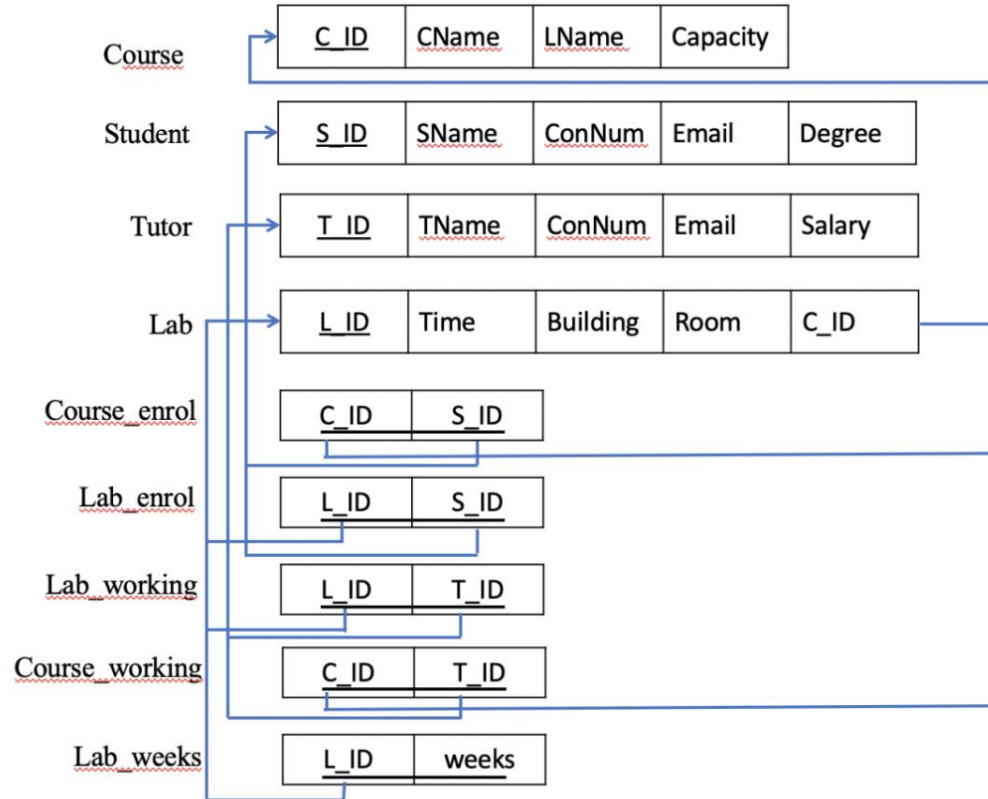
- 2.5 hrs., 9:30am – 12:00pm 6 Dec
- **Consultation: 25 Nov and 27 Nov 3-4pm**
- You are permitted ONE Double Sided Cheat Sheet (print/non-print)
- For the additional 30 minutes: I will add some conceptual questions
E.g. “Briefly explain the how the requirements of 3NF enforce the 1NF and the 2NF (hint: you may frame your discussion by the types of functional dependencies 3NF disallows) (4 marks)”
- That need you to be somewhat familiar with the meaning of the definitions

Remaining Matters of Business

- This week:
 - Me submit final exam to RES
 - Release of Assignment One Marks
 - Release of In-class assessment Solutions
 - Course Teaching Evaluation by Tomorrow
 - Some students with current grouping difficulties..
- This weekend:
 - Some students who have lost Oracle Logins
 - Some standard notations to expect/follow in Exam

If I ask you to draw a “relational schema”

- Draw **relational diagram** in exam as shown below



Project - Oracle Database

- *Select table_name from user_tables;*
- *describe <tablename>*

Connecting to Oracle Database

```
1. import java.sql.*;
2. Class.forName ("oracle.jdbc.driver.OracleDriver");
3. DriverManager.registerDriver (new
    oracle.jdbc.driver.OracleDriver());
4. Connection conn =
    DriverManager.getConnection("jdbc:oracle:thin:@//db18.cse.
    cuhk.edu.hk:1521/oradb.cse.cuhk.edu.hk", "<username>",
    "<password>");
5. Statement stmt = conn.createStatement();
6. String query =
7. stmt.execute(query);
```

Some Plans for Next Years

- Scrapping Group-based Project
 - 5 Assignments +
 - At least two coding: implementation of B+ Tree/ Buffer Pool in any pre-approved language of your choosing (C++, Python, Java)
- Participation (
 - Per week N: rank a given list of concepts in terms of understanding + a 100-200 statement on the why you don't understand your least understood concept
- A bigger Database (that CSE servers could hold..)
- Tutorials -> Lab Format?
 - Move all of SQL to lab content, and cover advanced SQL in lectures
 - Existing material -> Practice
- Welcome More Comments... put them in the CTE

Course and Teaching Evaluation (CTE) for Term 1, 2024-25

Notes for Students:

- The CTE links were sent to your **CUHK Link email accounts**. Please check.
- You will receive separate emails for teacher's evaluation and TA's evaluation.
- Subject of the email is "Online CTE Questionnaire for [course code]: [course name] – [Teacher or TA name]".
- No login is required and all feedback is anonymous.
- Course teacher will give you ~10-15 minutes to fill out the CTE questionnaires before the end of the class.
- **Submission deadline: By 23:59 tomorrow**